Local RAG Architecture Overview: Retrieval-Augmented Generation (RAG) enhances language models by

Vector Search and Embeddings: Documents are converted into fixed-dimensional dense vector embeddin

Foundry Local Inference: Microsoft Foundry Local provides high-performance local AI model execution on